

```
In [ ]: %load_ext autoreload
%autoreload 2

In [ ]: import open_clip
import torch
from sklearn.svm import LinearSVC
import numpy as np
from datasets import load_from_disk
from tqdm import tqdm
import sys
sys.path.append("../tools")
from utils import load_data_split
from sklearn.model_selection import cross_val_score
import warnings
import pandas as pd

In [ ]: model, _, preprocess = open_clip.create_model_and_transforms('hf-hub:laion/CLIP-ViT-L-14-DataComp.XL-s13B-b90K', device="cuda")

model.eval()
```

```
Out[ ]: CLIP(
  (visual): VisionTransformer(
    (conv1): Conv2d(3, 1024, kernel_size=(14, 14), stride=(14, 14), bias=False)
    (patch_dropout): Identity()
    (ln_pre): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
    (transformer): Transformer(
      (resblocks): ModuleList(
        (0-23): 24 x ResidualAttentionBlock(
          (ln_1): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
          (attn): MultiheadAttention(
            (out_proj): NonDynamicallyQuantizableLinear(in_features=1024, out_features=1024, bias=True)
          )
          (ls_1): Identity()
          (ln_2): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
          (mlp): Sequential(
            (c_fc): Linear(in_features=1024, out_features=4096, bias=True)
            (gelu): GELU(approximate='none')
            (c_proj): Linear(in_features=4096, out_features=1024, bias=True)
          )
          (ls_2): Identity()
        )
      )
    )
    (ln_post): LayerNorm((1024,), eps=1e-05, elementwise_affine=True)
  )
  (transformer): Transformer(
    (resblocks): ModuleList(
      (0-11): 12 x ResidualAttentionBlock(
        (ln_1): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
        (attn): MultiheadAttention(
          (out_proj): NonDynamicallyQuantizableLinear(in_features=768, out_features=768, bias=True)
        )
        (ls_1): Identity()
        (ln_2): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
        (mlp): Sequential(
          (c_fc): Linear(in_features=768, out_features=3072, bias=True)
          (gelu): GELU(approximate='none')
          (c_proj): Linear(in_features=3072, out_features=768, bias=True)
        )
        (ls_2): Identity()
      )
    )
  )
  (token_embedding): Embedding(49408, 768)
  (ln_final): LayerNorm((768,), eps=1e-05, elementwise_affine=True)
)
```

## Loading the local version of ELSA

```
In [ ]: dataset_path = "../data/big/"

In [ ]: ds_train = load_from_disk(dataset_path+"train")
```

## SVM Binary Classification

### Feature extraction with CLIP

```
In [ ]: def preprocess_img(X):
    with torch.no_grad():
        for i, img in enumerate(X):
            X[i] = model.encode_image(preprocess(img).unsqueeze(0).cuda()).detach().cpu().numpy()
    return X

In [ ]: X_train, y_train = preprocess_img(ds_train["image"]), ds_train["label"]

/home/lsaland/micromamba/envs/clip/lib/python3.11/site-packages/PIL/Image.py:1000: UserWarning: Palette images with Transparency expressed in bytes should be converted to RGBA images
  warnings.warn(

In [ ]: X_train = np.array([x.flatten() for x in X_train])

In [ ]: clf = LinearSVC()

clf.fit(X_train, y_train)

/home/lsaland/micromamba/envs/clip/lib/python3.11/site-packages/sklearn/svm/_classes.py:31: FutureWarning: The default value of `dual` will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly to suppress the warning.
  warnings.warn(
/home/lsaland/micromamba/envs/clip/lib/python3.11/site-packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
  warnings.warn(

Out[ ]: • LinearSVC ⓘ ?
LinearSVC()

In [ ]: ds_test = load_from_disk(dataset_path+"test")

In [ ]: X_test, y_test = preprocess_img(ds_test["image"]), ds_test["label"]

/home/lsaland/micromamba/envs/clip/lib/python3.11/site-packages/PIL/Image.py:1000: UserWarning: Palette images with Transparency expressed in bytes should be converted to RGBA images
  warnings.warn(

In [ ]: X_test = np.array([x.flatten() for x in X_test])
```

```
In [ ]: mean_accuracy_test = clf.score(X_test,y_test)
print("Mean accuracy on test set: ",mean_accuracy_test)

Mean accuracy on test set:  0.9819
```

```
In [ ]: X = np.vstack((X_train,X_test))
y = np.hstack((y_train,y_test))
```

```
In [ ]: warnings.filterwarnings("ignore", category=Warning)

cv = cross_val_score(clf,X,y,cv=10)
print("Cross validation score: ",np.mean(cv))

Cross validation score:  0.9846
```

## SVM on JPEG QF 40 data with training on any img format (training set is not JPEG only)

```
In [ ]: X_train, y_train = load_data_split("../data/big",
                                          split="train",
                                          model=model,
                                          preprocess=preprocess,
                                          device="cuda")
```

```
In [ ]: X_test, y_test = load_data_split("../data/big_QF_40",
                                          split="test",
                                          model=model,
                                          preprocess=preprocess,device="cuda")
```

```
In [ ]: clf = LinearSVC()
clf.fit(X_train,y_train)
clf.score(X_test,y_test)
```

```
Out[ ]: 0.7381
```

## SVM on JPEG QF 40 data with training on JPEG QF 40

```
In [ ]: X_train, y_train = load_data_split("../data/big_QF_40",
                                          split="train",
                                          model=model,
                                          preprocess=preprocess,
                                          device="cuda")
```

```
In [ ]: clf.fit(X_train,y_train)
clf.score(X_test,y_test)
```

```
Out[ ]: 0.9811
```

## SVM trained on JPEG QF 40 and tested on ELSA

```
In [ ]: X_test, y_test = load_data_split("../data/big",
                                          split="test",
                                          model=model,
                                          preprocess=preprocess,
                                          device="cuda")
```

/home/lsaland/micromamba/envs/clip/lib/python3.11/site-packages/PIL/Image.py:1000: UserWarning: Palette images with Transparency expressed in bytes should be converted to RGBA images  
warnings.warn(

```
In [ ]: clf.score(X_test,y_test)
```

```
Out[ ]: 0.9515
```

## SVM performance on different training size (train jpeg 40, test ELSA 10 000)

### Train size: 100

```
In [ ]: X_train_100, y_train_100 = load_data_split("../data/small_QF_40",
                                                    split="train",
                                                    model=model,
                                                    preprocess=preprocess,
                                                    device="cuda")
```

```
In [ ]: warnings.filterwarnings("ignore",category=FutureWarning)

clf = LinearSVC()
print("mean accuracy for train set of size 100: ",clf.fit(X_train_100,y_train_100).score(X_test,y_test))

mean accuracy for train set of size 100:  0.8805
```

### Train size: 1000

```
In [ ]: X_train_1000, y_train_1000 = load_data_split("../data/medium_QF_40",
                                                       split="train",
                                                       model=model,
                                                       preprocess=preprocess,
                                                       device="cuda")
```

```
In [ ]: from sklearn.exceptions import ConvergenceWarning
warnings.filterwarnings("ignore",category=ConvergenceWarning)

clf = LinearSVC()
print("mean accuracy for train set of size 1000: ",clf.fit(X_train_1000,y_train_1000).score(X_test,y_test))

mean accuracy for train set of size 1000:  0.9391
```

### Train size: 10 000

```
In [ ]: X_train_10000, y_train_10000 = load_data_split("../data/big_QF_40",
                                                         split="train",
                                                         model=model,
                                                         preprocess=preprocess,
                                                         device="cuda")
```

```
In [ ]: from sklearn.exceptions import ConvergenceWarning
warnings.filterwarnings("ignore",category=ConvergenceWarning)

clf = LinearSVC()
print("mean accuracy for train set of size 10 000: ",clf.fit(X_train_10000,y_train_10000).score(X_test,y_test))

mean accuracy for train set of size 10 000:  0.9513
```

## Training and testing for various quality factors

```
In [ ]: results = np.zeros((3,3))
k = 0

load_dataset_q = lambda q, split : load_data_split("../data/big_QF_" + str(q),
                                                    split=split,
                                                    model=model,
                                                    preprocess=preprocess,
                                                    device="cuda",
                                                    show_progress_bar=True)

train_data = [load_dataset_q(40,"train"),
              load_dataset_q(65,"train"),
              load_dataset_q(90,"train")]

test_data = [load_dataset_q(40,"test"),
             load_dataset_q(65,"test"),
             load_dataset_q(90,"test")]

100%|██████████| 10000/10000 [02:58<00:00, 56.12it/s]
100%|██████████| 10000/10000 [02:58<00:00, 55.96it/s]
100%|██████████| 10000/10000 [02:58<00:00, 56.01it/s]
100%|██████████| 10000/10000 [02:58<00:00, 56.12it/s]
100%|██████████| 10000/10000 [02:58<00:00, 56.12it/s]
100%|██████████| 10000/10000 [02:58<00:00, 56.05it/s]
```

```
In [ ]: warnings.filterwarnings("ignore",category=Warning)
clf = LinearSVC()
k = 0
with tqdm(total=9) as bar:
    for i in range(3):
        X_train, y_train = train_data[i]
        for j in range(3):
            X_test, y_test = test_data[j]
            results[i,j] = clf.fit(X_train,y_train).score(X_test,y_test)

            k += 1
            bar.n = k
            bar.refresh()

0%|          | 0/9 [00:00<?, ?it/s]100%|██████████| 9/9 [00:06<00:00, 1.47it/s]
```

```
In [ ]: results
```

```
Out[ ]: array([[0.9813, 0.973 , 0.9797],
               [0.945 , 0.9825, 0.983 ],
               [0.7744, 0.8581, 0.9925]])
```

```
In [ ]: qf = (40,65,90)
d = {"q_train" : [], "q_test": [], "accuracy": []}

for i in range(3):
    for j in range(3):
        d["q_train"].append(qf[i])
        d["q_test"].append(qf[j])
        d["accuracy"].append(results[i,j])
```

```
In [ ]: pd.DataFrame.from_records(results).to_csv("../docs/CLIP_jpg.csv")
```

```
Out[ ]:
   0    1    2
0 0.9813 0.9730 0.9797
1 0.9450 0.9825 0.9830
2 0.7744 0.8581 0.9925
```