

SVM classifier trained on images (real and generated ones from 4 generators) from ELSA compressed in jpg with quality 40 and coming. The classifier is tested on images from synthbuster

```
In [ ]: %load_ext autoreload
        %autoreload 2
```

```
In [ ]: import sys
        sys.path.append("../tools")
        from utils import load_data_features, load_synthbuster_balanced, map_synthbuster_classes
        from sklearn.svm import LinearSVC
        from datasets import load_from_disk
        import open_clip
        from IPython.display import clear_output
        import numpy as np
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
        import matplotlib.pyplot as plt
```

```
In [ ]: model, _, preprocess = open_clip.create_model_and_transforms('hf-hub:laion/CLIP-ViT-L-14-DataComp.XL-s13B-b90K', device="cuda")

        model.eval()
        clear_output()
```

```
In [ ]: def label_conversion(e):
        e["label"] = 1 if e["label"] == "real" else 0
        return e
```

```
In [ ]: X_sb, y_sb = load_synthbuster_balanced("../data/synthbuster_test",
        balance_real_fake=True,
        binary_classification=True)
```

```
In [ ]: X_40, y_40 = load_data_features("../data/big_QF_40_features", "train")
        X_h, y_h = load_data_features("../data/holistic", "train")
```

Train on big_QF_40_features (1 generator) / Test on holistic (4 generators)

```
In [ ]: clf = LinearSVC(dual="auto")
        clf.fit(X_40, y_40).score(X_h, y_h)
```

```
Out[ ]: 0.641
```

Train on holistic (4 generators) / Test on big_QF_40_features (1 generator)

```
In [ ]: clf = LinearSVC(dual="auto")
        clf.fit(X_h, y_h).score(X_40, y_40)
```

```
Out[ ]: 0.9426
```

Train on holistic (4 generators) / Test on synthbuster

```
In [ ]: clf.fit(X_h, y_h).score(X_sb, y_sb)
```

```
Out[ ]: 0.6444199116820554
```

Train on synthbuster (9 generators) / Test on holistic (4 generators)

```
In [ ]: clf.fit(X_sb, y_sb).score(X_h, y_h)
```

```
Out[ ]: 0.6202
```

Train and test on synthbuster

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X_sb, y_sb, shuffle=True, test_size=.2)
        print(X_train.shape)
        clf.fit(X_train, y_train).score(X_test, y_test)
```

```
(7971, 768)
Out[ ]: 0.9839438033115906
```

Train and test on big_QF_40 (1 generator)

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X_40, y_40, shuffle=True, test_size=.2)
        clf.fit(X_train, y_train).score(X_test, y_test)
```

```
Out[ ]: 0.981
```

Train on concatenation of synthbuster + holistic and test on holistic

train on synthbuster + holistic

```
In [ ]: X_sb_train, X_sb_test, y_sb_train, y_sb_test = train_test_split(X_sb, y_sb, test_size=.2, shuffle=True, random_state=7)
        X_h_train, y_h_train = load_data_features("../data/holistic", split="train")
        X_h_test, y_h_test = load_data_features("../data/holistic", split="test")

        X_train = np.vstack((X_sb_train, X_h_train))
        y_train = np.hstack((y_sb_train, y_h_train))

        clf.fit(X_train, y_train)
```

```
Out[ ]: LinearSVC ⓘ ?
        LinearSVC(dual='auto')
```

test on holistic

```
In [ ]: clf.score(X_h_test, y_h_test)
```

```
Out[ ]: 0.9508
```

test on synthbuster

```
In [ ]: clf.score(X_sb_test, y_sb_test)
```

```
Out[ ]: 0.9668840943301555
```

test on big_QF_40

```
In [ ]: X_40_test, y_40_test = load_data_features("../data/big_QF_40_features", split="test")
        clf.score(X_40, y_40)
```

```
Out[ ]: 0.9449
```

Multi class classifier

uniform distribution of data across all classes

```
In [ ]: from sklearn.multiclass import OneVsOneClassifier, OneVsRestClassifier

        ovo = OneVsOneClassifier(LinearSVC(dual="auto"))
        ovr = OneVsRestClassifier(LinearSVC(dual="auto"))
        X_sb, y_sb = load_synthbuster_balanced("../data/synthbuster_test",
                                              binary_classification=False,
                                              balance_real_fake=False)
```

```
In [ ]: from collections import Counter

        X_train, X_test, y_train, y_test = train_test_split(X_sb,
                                                            y_sb,
                                                            test_size=.2,
                                                            shuffle=True,
                                                            random_state=7)

        display(Counter(y_sb))
        display(Counter(y_test))
```

```
Counter({1: 1000,
          2: 1000,
          3: 1000,
          4: 1000,
          5: 1000,
          6: 1000,
          7: 1000,
          8: 1000,
          9: 1000,
          0: 1000})
Counter({4: 220,
          9: 215,
          8: 209,
          5: 206,
          3: 205,
          6: 200,
          2: 192,
          7: 189,
          0: 183,
          1: 181})
```

1000 points per classes

```
In [ ]: ovo.fit(X_train, y_train).score(X_test, y_test)
```

```
Out[ ]: 0.802
```

```
In [ ]: ovr.fit(X_train, y_train).score(X_test, y_test)
```

```
Out[ ]: 0.793
```

50% real images and 50% generated (9 generators)

```
In [ ]: X_sb, y_sb = load_synthbuster_balanced("../data/synthbuster_test",
                                              binary_classification=False,
                                              balance_real_fake=True)

        X_train, X_test2, y_train, y_test2 = train_test_split(X_sb,
                                                            y_sb,
                                                            shuffle=True,
                                                            random_state=7,
                                                            test_size=.2)

        X_train_filtered = np.zeros_like(X_train)
        y_train_filtered = np.zeros_like(y_train)
        k = 0
        for i, x in enumerate(X_train):
            if x not in X_test2:
                X_train_filtered[k] = X_train[i]
                y_train_filtered[k] = y_train[i]
                k += 1
        X_train_filtered = X_train_filtered[:k]
        y_train_filtered = y_train_filtered[:k]

        for x in X_train_filtered:
            if x in X_test2:
                print("PROBLEM: SOME POINT ARE SHARED BETWEEN TRAIN AND TEST")
        Counter(y_train_filtered)
```

```
Out[ ]: Counter({0: 3730,
          2: 369,
          1: 368,
          9: 355,
          7: 347,
          3: 347,
          8: 338,
          6: 336,
          4: 334,
          5: 330})
```

```
In [ ]: Counter(y_test2)
```

```
Out[ ]: Counter({4: 220,
          9: 215,
          8: 209,
          5: 206,
          3: 205,
          6: 200,
          2: 192,
          7: 189,
          0: 183,
          1: 181})
```

```
In [ ]: ovo.fit(X_train_filtered, y_train_filtered).score(X_test2, y_test2)
```

```
/home/lisaland/micromamba/envs/clip/lib/python3.11/site-packages/sklearn/svm/_base.py:1237: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(
0.7855
```

Out[]:

```
In [ ]: ovr.fit(X_train_filtered,y_train_filtered).score(X_test,y_test)
```

Out[]: 0.7465

Binary classification with multiclass classifier

```
In [ ]: y_test_hat_ovo = ovo.predict(X_test2)
result = np.zeros_like(y_test)
for i, y in enumerate(y_test_hat_ovo):
    if y == 0 and y_test2[i] == 0: # predicted real and was real
        result[i] = 1
    elif y > 0 and y_test2[i] > 0: # predicted fake and was really fake
        result[i] = 1
    else:
        result[i] = 0
np.mean(result)
```


Out[]: 0.987

Confusion matrix for binary classification

```
In [ ]: y_test_map = np.zeros_like(y_test2)
y_test_hat_map = np.zeros_like(y_test2)

for i in range(len(y_test2)):
    if y_test2[i] == 0:
        y_test_map[i] = 1
    if y_test_hat_ovo[i] == 0:
        y_test_hat_map[i] = 1
cm = confusion_matrix(y_test_map,y_test_hat_map)
ConfusionMatrixDisplay(cm,display_labels=["fake","real"]).plot()
```

Out[]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x75fa896abf10>



Confusion matrix for multiclass classification

```
In [ ]: y_test_hat_ovo = ovo.predict(X_test)

cm = confusion_matrix(y_test,y_test_hat_ovo)
labels = ['null',
          'glide',
          'stable-diffusion-1-4',
          'midjourney-v5',
          'stable-diffusion-xl',
          'dalle3',
          'firefly',
          'stable-diffusion-2',
          'stable-diffusion-1-3',
          'dalle2']

display = ConfusionMatrixDisplay(cm,display_labels=labels)
display.plot()
```

Out[]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x75fa943e84d0>

